

Assignment 4

February 3, 2021

*You are currently looking at **version 1.1** of this notebook. To download notebooks and datafiles, as well as get help on Jupyter notebooks in the Coursera platform, visit the [Jupyter Notebook FAQ](#) course resource.*

In []:

0.1 Assignment 4 - Understanding and Predicting Property Maintenance Fines

This assignment is based on a data challenge from the Michigan Data Science Team ([MDST](#)).

The Michigan Data Science Team ([MDST](#)) and the Michigan Student Symposium for Interdisciplinary Statistical Sciences ([MSSISS](#)) have partnered with the City of Detroit to help solve one of the most pressing problems facing Detroit - blight. [Blight violations](#) are issued by the city to individuals who allow their properties to remain in a deteriorated condition. Every year, the city of Detroit issues millions of dollars in fines to residents and every year, many of these fines remain unpaid. Enforcing unpaid blight fines is a costly and tedious process, so the city wants to know: how can we increase blight ticket compliance?

The first step in answering this question is understanding when and why a resident might fail to comply with a blight ticket. This is where predictive modeling comes in. For this assignment, your task is to predict whether a given blight ticket will be paid on time.

All data for this assignment has been provided to us through the [Detroit Open Data Portal](#). **Only the data already included in your Coursera directory can be used for training the model for this assignment.** Nonetheless, we encourage you to look into data from other Detroit datasets to help inform feature creation and model selection. We recommend taking a look at the following related datasets:

- [Building Permits](#)
- [Trades Permits](#)
- [Improve Detroit: Submitted Issues](#)
- [DPD: Citizen Complaints](#)
- [Parcel Map](#)

We provide you with two data files for use in training and validating your models: train.csv and test.csv. Each row in these two files corresponds to a single blight ticket, and includes information about when, why, and to whom each ticket was issued. The target variable is compliance,

which is True if the ticket was paid early, on time, or within one month of the hearing data, False if the ticket was paid after the hearing date or not at all, and Null if the violator was found not responsible. Compliance, as well as a handful of other variables that will not be available at test-time, are only included in train.csv.

Note: All tickets where the violators were found not responsible are not considered during evaluation. They are included in the training set as an additional source of data for visualization, and to enable unsupervised and semi-supervised approaches. However, they are not included in the test set.

File descriptions (Use only this data for training your model!)

readonly/train.csv - the training set (all tickets issued 2004-2011)

readonly/test.csv - the test set (all tickets issued 2012-2016)

readonly/addresses.csv & readonly/latlons.csv - mapping from ticket id to addresses

Note: misspelled addresses may be incorrectly geolocated.

Data fields

train.csv & test.csv

ticket_id - unique identifier for tickets

agency_name - Agency that issued the ticket

inspector_name - Name of inspector that issued the ticket

violation_name - Name of the person/organization that the ticket was issued to

violation_street_number, violation_street_name, violation_zip_code - Address where

mailing_address_str_number, mailing_address_str_name, city, state, zip_code, non_us

ticket_issued_date - Date and time the ticket was issued

hearing_date - Date and time the violator's hearing was scheduled

violation_code, violation_description - Type of violation

disposition - Judgment and judgement type

fine_amount - Violation fine amount, excluding fees

admin_fee - \$20 fee assigned to responsible judgments

state_fee - \$10 fee assigned to responsible judgments late_fee - 10% fee assigned to responsible

judgments discount_amount - discount applied, if any clean_up_cost - DPW clean-up or graffiti

removal cost judgment_amount - Sum of all fines and fees graffiti_status - Flag for graffiti viola-

tions

train.csv only

payment_amount - Amount paid, if any

payment_date - Date payment was made, if it was received

payment_status - Current payment status as of Feb 1 2017

balance_due - Fines and fees still owed

collection_status - Flag for payments in collections

compliance [target variable for prediction]

Null = Not responsible

0 = Responsible, non-compliant

1 = Responsible, compliant

compliance_detail - More information on why each ticket was marked compliant or non-

0.2 Evaluation

Your predictions will be given as the probability that the corresponding blight ticket will be paid on time.

The evaluation metric for this assignment is the Area Under the ROC Curve (AUC).

Your grade will be based on the AUC score computed for your classifier. A model which with an AUROC of 0.7 passes this assignment, over 0.75 will receive full points. ____

For this assignment, create a function that trains a model to predict blight ticket compliance in Detroit using `readonly/train.csv`. Using this model, return a series of length 61001 with the data being the probability that each corresponding ticket from `readonly/test.csv` will be paid, and the index being the `ticket_id`.

Example:

```
ticket_id
284932    0.531842
285362    0.401958
285361    0.105928
285338    0.018572
...
376499    0.208567
376500    0.818759
369851    0.018528
Name: compliance, dtype: float32
```

0.2.1 Hints

- Make sure your code is working before submitting it to the autograder.
- Print out your result to see whether there is anything weird (e.g., all probabilities are the same).
- Generally the total runtime should be less than 10 mins. You should NOT use Neural Network related classifiers (e.g., `MLPClassifier`) in this question.
- Try to avoid global variables. If you have other functions besides `blight_model`, you should move those functions inside the scope of `blight_model`.
- Refer to the pinned threads in Week 4's discussion forum when there is something you could not figure it out.

```
In [1]: import pandas as pd
        df = pd.read_csv('train.csv', encoding = "ISO-8859-1")

        df.compliance = df.compliance.fillna(value=-1)
        df = df[df.compliance != -1]

        df.head(10)
```

```
/opt/conda/lib/python3.6/site-packages/IPython/core/interactiveshell.py:2717: Dtype
interactivity=interactivity, compiler=compiler, result=result)
```

```
Out[1]:
```

	ticket_id	agency_name	\
0	22056	Buildings, Safety Engineering & Env Department	
1	27586	Buildings, Safety Engineering & Env Department	
5	22046	Buildings, Safety Engineering & Env Department	
6	18738	Buildings, Safety Engineering & Env Department	
7	18735	Buildings, Safety Engineering & Env Department	
8	18733	Buildings, Safety Engineering & Env Department	
9	28204	Buildings, Safety Engineering & Env Department	
12	18743	Buildings, Safety Engineering & Env Department	
13	18741	Buildings, Safety Engineering & Env Department	
14	18978	Buildings, Safety Engineering & Env Department	

	inspector_name	violator_name	\
0	Sims, Martinzie	INVESTMENT INC., MIDWEST MORTGAGE	
1	Williams, Darrin	Michigan, Covenant House	
5	Sims, Martinzie	KASIMU, UKWELI	
6	Williams, Darrin	Deerwood Development Group Inc, Deer	
7	Williams, Darrin	Rafee Auto Services L.L.C., RAF	
8	Williams, Darrin	Rafee Auto Services L.L.C., RAF	
9	Williams, Darrin	Inc, Nanno	
12	Williams, Darrin	Gardner Resale, GAR	
13	Williams, Darrin	Hardaway, Kevin	
14	Williams, Darrin	TLC Hand Car Wash, a/k/a	

	violation_street_number	violation_street_name	violation_zip_code	\
0	2900.0	TYLER	NaN	
1	4311.0	CENTRAL	NaN	
5	6478.0	NORTHFIELD	NaN	
6	8027.0	BRENTWOOD	NaN	
7	8228.0	MT ELLIOTT	NaN	
8	8228.0	MT ELLIOTT	NaN	
9	15307.0	SEVEN MILE	NaN	
12	9100.0	VAN DYKE	NaN	
13	20024.0	SCHAEFER	NaN	
14	9425.0	VAN DYKE	NaN	

	mailing_address_str_number	mailing_address_str_name	city	\
0	3.0	S. WICKER	CHICAGO	
1	2959.0	Martin Luther King	Detroit	
5	2755.0	E. 17TH	LOG BEACH	
6	476.0	Garfield	Clinton	
7	8228.0	Mt. Elliott	Detroit	
8	8228.0	Mt. Elliott	Detroit	
9	1537.0	E. Seven Mile	Detroit	

12	91.0	Van Dyke	Detroit
13	224.0	Schaefer	Detroit
14	9425.0	Van Dyke	Detroit

	...	clean_up_cost	judgment_amount	payment_amount	balance_due	\
0	...	0.0	305.0	0.0	305.0	
1	...	0.0	855.0	780.0	75.0	
5	...	0.0	305.0	0.0	305.0	
6	...	0.0	855.0	0.0	855.0	
7	...	0.0	140.0	0.0	140.0	
8	...	0.0	140.0	0.0	140.0	
9	...	0.0	855.0	0.0	855.0	
12	...	0.0	855.0	0.0	855.0	
13	...	0.0	855.0	0.0	855.0	
14	...	0.0	855.0	0.0	855.0	

	payment_date	payment_status	collection_status	grafitti_stat
0	NaN	NO PAYMENT APPLIED	NaN	N
1	2005-06-02 00:00:00	PAID IN FULL	NaN	N
5	NaN	NO PAYMENT APPLIED	NaN	N
6	NaN	NO PAYMENT APPLIED	NaN	N
7	NaN	NO PAYMENT APPLIED	NaN	N
8	NaN	NO PAYMENT APPLIED	NaN	N
9	NaN	NO PAYMENT APPLIED	NaN	N
12	NaN	NO PAYMENT APPLIED	IN COLLECTION	N
13	NaN	NO PAYMENT APPLIED	IN COLLECTION	N
14	NaN	NO PAYMENT APPLIED	NaN	N

	compliance_detail	compliance
0	non-compliant by no payment	0.0
1	compliant by late payment within 1 month	1.0
5	non-compliant by no payment	0.0
6	non-compliant by no payment	0.0
7	non-compliant by no payment	0.0
8	non-compliant by no payment	0.0
9	non-compliant by no payment	0.0
12	non-compliant by no payment	0.0
13	non-compliant by no payment	0.0
14	non-compliant by no payment	0.0

[10 rows x 34 columns]

```
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import roc_auc_score
from sklearn.metrics import roc_curve
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import GridSearchCV
```

```
def blight_model(): return ans
```

```
In [2]: import pandas as pd
```

```

import numpy as np
import math
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import roc_auc_score
from sklearn.metrics import roc_curve, auc
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import GridSearchCV

df = pd.read_csv('train.csv', encoding = "ISO-8859-1")
df.index = df['ticket_id']
features_name = ['fine_amount', 'admin_fee', 'state_fee', 'late_fee']

df.compliance = df.compliance.fillna(value=-1)
df = df[df.compliance != -1]

X = df[features_name]
X.fillna(value = -1)

y = df.compliance

X_train, X_test, y_train, y_test = train_test_split(X, y, random_state = 0)

clf = RandomForestClassifier(n_estimators = 10, max_depth = 5).fit(X_train,
features_name = ['fine_amount', 'admin_fee', 'state_fee', 'late_fee']

df_test = pd.read_csv('readonly/test.csv', encoding = "ISO-8859-1")

df_test.index = df_test['ticket_id']

X_predict = clf.predict_proba(df_test[features_name])

ans = pd.Series(data = X_predict[:,1], index = df_test['ticket_id'], dtype=
#return ans

```

```

/opt/conda/lib/python3.6/site-packages/IPython/core/interactiveshell.py:2717: Dtype
interactivity=interactivity, compiler=compiler, result=result)

```

```

In [3]: df.compliance = df.compliance.fillna(value=-1)
df = df[df.compliance != -1]

```

```

In [5]: clf.score(X_test, y_test)

```

```

Out[5]: 0.92847135351513632

```

```

In [6]: import pandas as pd
import numpy as np

```

```
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
```

```
df = pd.read_csv('train.csv', encoding = "ISO-8859-1")
df_test = pd.read_csv('readonly/test.csv', encoding = "ISO-8859-1")
```

```
/opt/conda/lib/python3.6/site-packages/IPython/core/interactiveshell.py:2717: Dtype
interactivity=interactivity, compiler=compiler, result=result)
```

```
In [7]: df.head(4)
```

```
Out[7]:
```

	ticket_id	agency_name	
0	22056	Buildings, Safety Engineering & Env Department	
1	27586	Buildings, Safety Engineering & Env Department	
2	22062	Buildings, Safety Engineering & Env Department	
3	22084	Buildings, Safety Engineering & Env Department	

	inspector_name	violator_name	
0	Sims, Martinzie	INVESTMENT INC., MIDWEST MORTGAGE	
1	Williams, Darrin	Michigan, Covenant House	
2	Sims, Martinzie	SANDERS, DERRON	
3	Sims, Martinzie	MOROSI, MIKE	

	violation_street_number	violation_street_name	violation_zip_code	
0	2900.0	TYLER	NaN	
1	4311.0	CENTRAL	NaN	
2	1449.0	LONGFELLOW	NaN	
3	1441.0	LONGFELLOW	NaN	

	mailing_address_str_number	mailing_address_str_name	city	...
0	3.0	S. WICKER	CHICAGO	...
1	2959.0	Martin Luther King	Detroit	...
2	23658.0	P.O. BOX	DETROIT	...
3	5.0	ST. CLAIR	DETROIT	...

	clean_up_cost	judgment_amount	payment_amount	balance_due	
0	0.0	305.0	0.0	305.0	
1	0.0	855.0	780.0	75.0	
2	0.0	0.0	0.0	0.0	
3	0.0	0.0	0.0	0.0	

	payment_date	payment_status	collection_status	grafitti_status	
0	NaN	NO PAYMENT APPLIED	NaN	NaN	
1	2005-06-02 00:00:00	PAID IN FULL	NaN	NaN	
2	NaN	NO PAYMENT APPLIED	NaN	NaN	
3	NaN	NO PAYMENT APPLIED	NaN	NaN	

	compliance_detail	compliance
0	non-compliant by no payment	0.0
1	compliant by late payment within 1 month	1.0
2	not responsible by disposition	NaN
3	not responsible by disposition	NaN

[4 rows x 34 columns]

In [8]: df_test.head(4)

```
Out[8]:
```

	ticket_id	agency_name	inspector_name	\
0	284932	Department of Public Works	Granberry, Aisha B	
1	285362	Department of Public Works	Lusk, Gertrina	
2	285361	Department of Public Works	Lusk, Gertrina	
3	285338	Department of Public Works	Talbert, Reginald	

	violator_name	violation_street_number	violation_street_name	\
0	FLUELLEN, JOHN A	10041.0	ROSEBERRY	
1	WHIGHAM, THELMA	18520.0	EVERGREEN	
2	WHIGHAM, THELMA	18520.0	EVERGREEN	
3	HARABEDIEN, POPKIN	1835.0	CENTRAL	

	violation_zip_code	mailing_address_str_number	mailing_address_str_name	\
0	NaN	141	ROSEBERRY	
1	NaN	19136	GLASTONBURY	
2	NaN	19136	GLASTONBURY	
3	NaN	2246	NELSON	

	city	...	\
0	DETROIT	...	
1	DETROIT	...	
2	DETROIT	...	
3	WOODHAVEN	...	

	violation_description	dispositio
0	Failure to secure City or Private solid waste ...	Responsible by Defaul
1	Allowing bulk solid waste to lie or accumulate...	Responsible by Defaul
2	Improper placement of Courville container betw...	Responsible by Defaul
3	Allowing bulk solid waste to lie or accumulate...	Responsible by Defaul

	fine_amount	admin_fee	state_fee	late_fee	discount_amount	clean_up_cost
0	200.0	20.0	10.0	20.0	0.0	0.0
1	1000.0	20.0	10.0	100.0	0.0	0.0
2	100.0	20.0	10.0	10.0	0.0	0.0
3	200.0	20.0	10.0	20.0	0.0	0.0

	judgment_amount	grafitti_status
--	-----------------	-----------------

0	250.0	NaN
1	1130.0	NaN
2	140.0	NaN
3	250.0	NaN

[4 rows x 27 columns]

In [9]: df.shape

Out[9]: (250306, 34)

In [10]: df_test.shape

Out[10]: (61001, 27)

In [11]: df_test.describe()

Out[11]:

	ticket_id	violation_street_number	non_us_str_code	fine_amount
count	61001.000000	6.100100e+04	0.0	61001.000000
mean	331724.532811	1.256638e+04	NaN	272.7141
std	25434.932141	1.414373e+05	NaN	360.1018
min	284932.000000	-1.512600e+04	NaN	0.0000
25%	310111.000000	6.008000e+03	NaN	50.0000
50%	332251.000000	1.213400e+04	NaN	200.0000
75%	353031.000000	1.716500e+04	NaN	250.0000
max	376698.000000	2.010611e+07	NaN	10000.0000

	admin_fee	state_fee	late_fee	discount_amount	clean_up_cost
count	61001.0	61001.0	61001.000000	61001.000000	61001.000000
mean	20.0	10.0	25.116219	0.239340	20.649711
std	0.0	0.0	36.310155	3.245894	242.375180
min	20.0	10.0	0.000000	0.000000	0.000000
25%	20.0	10.0	5.000000	0.000000	0.000000
50%	20.0	10.0	10.000000	0.000000	0.000000
75%	20.0	10.0	25.000000	0.000000	0.000000
max	20.0	10.0	1000.000000	250.000000	15309.000000

	judgment_amount
count	61001.000000
mean	347.895541
std	460.058043
min	0.000000
25%	85.000000
50%	250.000000
75%	305.000000
max	15558.800000


In [15]: import matplotlib.pyplot as plt

graph = df.iloc[:, :26]

graph.plot.scatter(x=0, y=1, c='DarkBlue', figsize=(16, 8))

plt.show

Out[15]: <function matplotlib.pyplot.show>



```
In [20]: import pandas as pd
import numpy as np
import math
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import roc_auc_score
from sklearn.metrics import roc_curve, auc
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import GridSearchCV

def blight_model():

    # Your code here

    df = pd.read_csv('train.csv', encoding = "ISO-8859-1")

    df.index = df['ticket_id']

    features_name = ['fine_amount', 'admin_fee', 'state_fee', 'late_fee']

    df.compliance = df.compliance.fillna(value=-1)

    df = df[df.compliance != -1]

    X = df[features_name]

    X.fillna(value = -1)
```

```

y = df.compliance

X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=42)

clf = RandomForestClassifier(n_estimators = 10, max_depth = 5).fit(X_train, y_train)

features_name = ['fine_amount', 'admin_fee', 'state_fee', 'late_fee']

df_test = pd.read_csv('readonly/test.csv', encoding = "ISO-8859-1")

df_test.index = df_test['ticket_id']

X_predict = clf.predict_proba(df_test[features_name])
predict = list(X_predict)
ans = pd.Series(data = X_predict[:,1], index = df_test['ticket_id'], dtype = float)

return ans

```

In [21]: blight_model()

/opt/conda/lib/python3.6/site-packages/IPython/core/interactiveshell.py:2827: DtypeWarning: Columns (0) have mixed types. Specify dtype option on pandas DataFrame constructors.
if self.run_code(code, result):

```

Out[21]: ticket_id
284932      0.060381
285362      0.026680
285361      0.069929
285338      0.060381
285346      0.069929
285345      0.060381
285347      0.057021
285342      0.402546
285530      0.026680
284989      0.026680
285344      0.057021
285343      0.026680
285340      0.026680
285341      0.057021
285349      0.069929
285348      0.060381
284991      0.026680
285532      0.026680
285406      0.026680
285001      0.026680
285006      0.026680
285405      0.026680

```

```

285337    0.026680
285496    0.057021
285497    0.060381
285378    0.026680
285589    0.026680
285585    0.060381
285501    0.069929
285581    0.026680
...
376367    0.026680
376366    0.036155
376362    0.036155
376363    0.060381
376365    0.026680
376364    0.036155
376228    0.036155
376265    0.036155
376286    0.367978
376320    0.036155
376314    0.036155
376327    0.367978
376385    0.367978
376435    0.465922
376370    0.367978
376434    0.057021
376459    0.069929
376478    0.005052
376473    0.036155
376484    0.024648
376482    0.026680
376480    0.026680
376479    0.026680
376481    0.026680
376483    0.036155
376496    0.026680
376497    0.026680
376499    0.069929
376500    0.069929
369851    0.303633
dtype: float32


```

```
In [16]: import matplotlib.pyplot as plt
```

```

graph=df.iloc[:,0::33]
graph.plot.scatter(x=0,y=1,c='DarkBlue',figsize=(16,8),s=1)
plt.show()

```



In []: